

Zhaoxuan Tan

Ph.D. Candidate, Computer Science and Engineering, University of Notre Dame

Email: ztan3@nd.edu / tanzx9@gmail.com Homepage: zhaoxuan.info

[Google Scholar](#) · [GitHub](#) · [LinkedIn](#)

Summary

Final-year Ph.D. candidate researching **LLM personalization and collaborative agents**, with a focus on user modeling, preference alignment, and efficient adaptation. **Five research internships** at **Google, Amazon, and Microsoft**; 40+ research papers, including **10 first-author papers**, and **3,000+ citations** (h-index: 24). Seeking full-time **Research Scientist / Applied Scientist** roles; **earliest start date: January 2027**.

Education

University of Notre Dame, Notre Dame, IN Aug 2023 – May 2027 (expected)

Ph.D. in Computer Science and Engineering; advisor: Prof. [Meng Jiang](#)

Xi'an Jiaotong University, Xi'an, China Aug 2019 – Jul 2023

B.E. in Computer Science and Technology; advisor: Prof. [Minnan Luo](#)

Industry Research Experience

Microsoft, Office of Applied Research, Redmond, WA May 2026 – Present

Research Intern; hosts: Dr. [Pei Zhou](#), Dr. [Mengting Wan](#)

- Developing **shared-context briefs for proactive agents** in long-term, multi-person Teams collaboration, distinguishing decisions from proposals while preserving attribution, unresolved disagreements, and evolving project state. A version of the briefer is already **in use internally**.
- Evaluating **claim grounding, attribution, and freshness** using GroupMemBench-based simulations with oracle references; early results show stronger recovery of past state than prediction of future opportunities and risks.

Google, CoreML, Remote Oct 2025 – Apr 2026

Student Researcher; hosts: Dr. [Yan Zhu](#), Dr. [Xiang Zhai](#)

- Developed **BLUE**, a reinforcement learning framework that trains an LLM to generate interpretable user profiles using recommendation rewards and a next-item prediction objective.
- Improved recommendation across **five domains** with a frozen embedding model, including zero-shot transfer from Amazon Clothing to three Amazon categories and Google Local Review. First-author paper: **NeurIPS 2026**.

Amazon, Rufus, Palo Alto, CA Jun 2025 – Oct 2025

Applied Scientist Intern; hosts: Dr. [Zixuan Zhang](#), Dr. [Haoyang Wen](#)

- Developed **Profile-to-PEFT**, an end-to-end hypernetwork that maps user profiles directly to LoRA adapters, eliminating per-user training at deployment and generalizing to unseen users.
- Outperformed prompt-based personalization and per-user fine-tuning (OPPU); generated personalized adapters in **0.57 s per user**, achieving **33× faster adaptation** than the fastest OPPU variant. First-author paper: **ACL 2026**.

Microsoft, Office of Applied Research, Remote Mar 2025 – Jun 2025

Research Intern (part-time); hosts: Dr. [Pei Zhou](#), Dr. [Mengting Wan](#)

- Built the **Meeting Dynamic Forecasting Benchmark (MDFB)** from **2,207 real-world meetings** and **24,794 forward-looking queries** to evaluate LLMs' ability to forecast multi-party discussions, tracking agendas, roles, and disagreement.
- Contributed to **GLARE**, an adversarial-reward post-training method achieving average human-evaluated win rates of **66% on utility** and **70% on human-likeness** across model and human baselines.

Amazon, Rufus, Palo Alto, CA May 2024 – Oct 2024

Applied Scientist Intern; hosts: Dr. [Zheng Li](#), Dr. [Tianyi Liu](#)

- Developed **PUGC**, converting unlabeled user-generated content into scalable preference data for DPO.
- Achieved a **35.93% length-controlled win rate on AlpacaEval 2.0** with Mistral-7B-Instruct-v0.2, **9.37 percentage points** above DPO with UltraFeedback preference data. First-author paper: **ACL 2025**.

Publications

* denotes equal contribution. Citation metrics and updates: [Google Scholar](#).

First-Author Publications

- [1] **Bridging Textual Profiles and Latent User Embeddings for Personalization.**
Zhaoxuan Tan, Xiang Zhai, Yan Zhu, Meng Jiang, Mohamed Hammad. *NeurIPS* 2026.
- [2] **Instant Personalized Large Language Model Adaptation via Hypernetwork.**
Zhaoxuan Tan, Zixuan Zhang, Haoyang Wen, Zheng Li, Rongzhi Zhang, Pei Chen, Fengran Mo, Zheyuan Liu, Qingkai Zeng, Qingyu Yin, Meng Jiang. *ACL* 2026.
- [3] **Aligning Large Language Models with Implicit Preferences from User-Generated Content.**
Zhaoxuan Tan, Zheng Li, Tianyi Liu, Haodong Wang, Hyokun Yun, Ming Zeng, Pei Chen, Zhihan Zhang, Yifan Gao, Ruijie Wang, Priyanka Nigam, Bing Yin, Meng Jiang. *ACL* 2025.
- [4] **Can Large Language Models Understand Preferences in Personalized Recommendation?**
Zhaoxuan Tan, Zinan Zeng, Qingkai Zeng, Zhenyu Wu, Zheyuan Liu, Fengran Mo, Meng Jiang. *arXiv preprint* 2025.
- [5] **Personalized Pieces: Efficient Personalized Large Language Models through Collaborative Efforts.**
Zhaoxuan Tan, Zheyuan Liu, Meng Jiang. *EMNLP* 2024.
- [6] **Democratizing Large Language Models via Personalized Parameter-Efficient Fine-tuning.**
Zhaoxuan Tan, Qingkai Zeng, Yijun Tian, Zheyuan Liu, Bing Yin, Meng Jiang. *EMNLP* 2024.
- [7] **User Modeling in the Era of Large Language Models: Current Research and Future Directions.**
Zhaoxuan Tan, Meng Jiang. *IEEE Data Engineering Bulletin* 2023.
- [8] **BotPercent: Estimating Bot Populations in Twitter Communities.**
Zhaoxuan Tan*, Shangbin Feng*, Melanie Sclar, Herun Wan, Minnan Luo, Yejin Choi, Yulia Tsvetkov. *Findings of EMNLP* 2023.
- [9] **KRACL: Contrastive Learning with Graph Context Modeling for Sparse Knowledge Graph Completion.**
Zhaoxuan Tan, Zilong Chen, Shangbin Feng, Qingyue Zhang, Qinghua Zheng, Jundong Li, Minnan Luo. *The Web Conference (WWW)* 2023.
- [10] **Twibot-22: Towards Graph-Based Twitter Bot Detection.**
Shangbin Feng*, Zhaoxuan Tan*, Herun Wan*, Ningnan Wang*, Zilong Chen*, Binchi Zhang*, . . . , Jundong Li, Minnan Luo. *NeurIPS 2022, Datasets and Benchmarks Track*.

Co-Authored Publications

- [11] **Prompt-Activation Duality: Improving Activation Steering via Attention-Level Interventions.**
Diancheng Kang, Zheyuan Liu, Ningshan Ma, Yue Huang, Zhaoxuan Tan, Meng Jiang. *NeurIPS* 2026.
- [12] **Steering Multimodal Large Language Models Decoding for Context-Aware Safety.**
Zheyuan Liu, Zhangchen Xu, Guangyao Dou, Xiangchi Yuan, Zhaoxuan Tan, Radha Poovendran, Meng Jiang. *EMNLP* 2026.
- [13] **DRIFT: Learning from Abundant User Dissatisfaction in Real-World Preference Learning.**
Yifan Wang, Bolian Li, Junlin Wu, Zhaoxuan Tan, Zheli Liu, Ruqi Zhang, Ananth Grama, Qingkai Zeng. *ICLR* 2026.
- [14] **From Personal to Collective: On the Role of Local and Global Memory in LLM Personalization.**
Zehong Wang, Junlin Wu, Zhaoxuan Tan, Bolian Li, Xianrui Zhong, Zheli Liu, Qingkai Zeng. *Findings of ACL* 2026.
- [15] **Agentic Conversational Search with Contextualized Reasoning via Reinforcement Learning.**
Fengran Mo, Yifan Gao, Sha Li, Hansi Zeng, Xin Liu, Zhaoxuan Tan, Xian Li, Jianshu Chen, Dakuo Wang, Meng Jiang. *Findings of ACL* 2026.
- [16] **MTMCS-Bench: Evaluating Contextual Safety of Multimodal Large Language Models in Multi-Turn Dialogues.**
Zheyuan Liu, Dongwhi Kim, Yixin Wan, Xiangchi Yuan, Zhaoxuan Tan, Fengran Mo, Meng Jiang. *Findings of ACL* 2026.
- [17] **Generalizable LLM Learning of Graph Synthetic Data with Post-training Alignment.**
Yizhuo Zhang*, Heng Wang*, Shangbin Feng*, Zhaoxuan Tan, Xinyun Liu, Yulia Tsvetkov. *Findings of ACL* 2026.
- [18] **GLARE: Generative Learning via Adversarial Reward Estimation for Social Dynamics Forecasting.**
Tenghao Huang, Zhaoxuan Tan, Muhao Chen, Jonathan May, Mengting Wan, Longqi Yang, Pei Zhou, Sihao Chen. *arXiv preprint* 2026.

- [19] **MoCo: A One-Stop Shop for Model Collaboration Research.**
Shangbin Feng*, Yuyang Bai*, Ziyuan Yang*, Yike Wang, **Zhaoxuan Tan**, Jiajie Yan, Zhenyu Lei, Wenxuan Ding, Weijia Shi, . . . , Luke Zettlemoyer, Yejin Choi, Yulia Tsvetkov. *arXiv preprint 2026*.
- [20] **Adaptive Personalized Conversational Information Retrieval.**
Fengran Mo, Yuchen Hui, Yuxing Tian, **Zhaoxuan Tan**, Chuan Meng, Zhan Su, Kaiyu Huang, Jian-Yun Nie. *CIKM 2025*.
- [21] **Enhancing Mathematical Reasoning in LLMs by Stepwise Correction.**
Zhenyu Wu, Qingkai Zeng, Zhihan Zhang, **Zhaoxuan Tan**, Chao Shen, Meng Jiang. *ACL 2025*.
- [22] **Modality-Aware Neuron Pruning for Unlearning in Multimodal Large Language Models.**
Zheyuan Liu, Guangyao Dou, Xiangchi Yuan, Chunhui Zhang, **Zhaoxuan Tan**, Meng Jiang. *ACL 2025*.
- [23] **CodeTaxo: Enhancing Taxonomy Expansion with Limited Examples via Code Language Prompts.**
Qingkai Zeng, Yuyang Bai, **Zhaoxuan Tan**, Zhenyu Wu, Shangbin Feng, Meng Jiang. *Findings of ACL 2025*.
- [24] **IHEval: Evaluating Language Models on Following the Instruction Hierarchy.**
Zhihan Zhang, Shiyang Li, Zixuan Zhang, Xin Liu, Haoming Jiang, Xianfeng Tang, Yifan Gao, Zheng Li, Haodong Wang, **Zhaoxuan Tan**, Yichuan Li, Qingyu Yin, Bing Yin, Meng Jiang. *NAACL 2025*.
- [25] **Protecting Privacy in Multimodal Large Language Models with MLLMU-Bench.**
Zheyuan Liu, Guangyao Dou, Mengzhao Jia, **Zhaoxuan Tan**, Qingkai Zeng, Yongle Yuan, Meng Jiang. *NAACL 2025*.
- [26] **Empirical Guidelines for Deploying LLMs onto Resource-constrained Edge Devices.**
Ruiyang Qin, Dancheng Liu, Zheyu Yan, **Zhaoxuan Tan**, Zixuan Pan, Zheng Jia, Meng Jiang, Ahmed Abbasi, Jinjun Xiong, Yiyu Shi. *ACM Transactions on Design Automation of Electronic Systems 2025*.
- [27] **KGQuiz: Evaluating the Generalization of Encoded Knowledge in Large Language Models.**
Yuyang Bai*, Shangbin Feng*, Vidhisha Balachandran, **Zhaoxuan Tan**, Shiqi Lou, Tianxing He, Yulia Tsvetkov. *The Web Conference (WWW) 2024, Oral*.
- [28] **Large Language Models Can Self-Correct with Key Condition Verification.**
Zhenyu Wu, Qingkai Zeng, Zhihan Zhang, **Zhaoxuan Tan**, Chao Shen, Meng Jiang. *EMNLP 2024*.
- [29] **Can LLM Graph Reasoning Generalize beyond Pattern Memorization?**
Yizhuo Zhang*, Heng Wang*, Shangbin Feng*, **Zhaoxuan Tan**, Xiaochuang Han, Tianxing He, Yulia Tsvetkov. *Findings of EMNLP 2024*.
- [30] **Chain-of-Layer: Iteratively Prompting Large Language Models for Taxonomy Induction from Limited Examples.**
Qingkai Zeng*, Yuyang Bai*, **Zhaoxuan Tan**, Shangbin Feng, Zhenwen Liang, Zhihan Zhang, Meng Jiang. *CIKM 2024*.
- [31] **What Does the Bot Say? Opportunities and Risks of Large Language Models in Social Media Bot Detection.**
Shangbin Feng, Herun Wan, Ningnan Wang, **Zhaoxuan Tan**, Minnan Luo, Yulia Tsvetkov. *ACL 2024*.
- [32] **DELL: Generating Reactions and Explanations for LLM-Based Misinformation Detection.**
Herun Wan*, Shangbin Feng*, **Zhaoxuan Tan**, Heng Wang, Yulia Tsvetkov, Minnan Luo. *Findings of ACL 2024*.
- [33] **Towards Safer Large Language Models through Machine Unlearning.**
Zheyuan Liu, Guangyao Dou, **Zhaoxuan Tan**, Yijun Tian, Meng Jiang. *Findings of ACL 2024*.
- [34] **Knowledge Crosswords: Geometric Knowledge Reasoning with Large Language Models.**
Wenxuan Ding*, Shangbin Feng*, Yuhao Liu, **Zhaoxuan Tan**, Vidhisha Balachandran, Tianxing He, Yulia Tsvetkov. *Findings of ACL 2024*.
- [35] **LMBot: Distilling Graph Knowledge into Language Model for Graph-less Deployment in Twitter Bot Detection.**
Zijian Cai, **Zhaoxuan Tan**, Zhenyu Lei, Zifeng Zhu, Hongrui Wang, Qinghua Zheng, Minnan Luo. *WSDM 2024*.
- [36] **Machine Unlearning in Generative AI: A Survey.**
Zheyuan Liu, Guangyao Dou, **Zhaoxuan Tan**, Yijun Tian, Meng Jiang. *arXiv preprint 2024*.
- [37] **Can Language Models Solve Graph Problems in Natural Language?**
Heng Wang, Shangbin Feng, Tianxing He, **Zhaoxuan Tan**, Xiaochuang Han, Yulia Tsvetkov. *NeurIPS 2023, Spotlight*.
- [38] **Detecting Spoilers in Movie Reviews with External Movie Knowledge and User Networks.**
Heng Wang, Wenqian Zhang, Yuyang Bai, **Zhaoxuan Tan**, Shangbin Feng, Qinghua Zheng, Minnan Luo. *EMNLP 2023*.

- [39] **KALM: Knowledge-Aware Integration of Local, Document, and Global Contexts for Long Document Understanding.**
Shangbin Feng, **Zhaoxuan Tan**, Wenqian Zhang, Zhenyu Lei, Yulia Tsvetkov. *ACL* 2023.
- [40] **BotMoE: Twitter Bot Detection with Community-Aware Mixtures of Modal-Specific Experts.**
Yuhan Liu, **Zhaoxuan Tan**, Heng Wang, Shangbin Feng, Qinghua Zheng, Minnan Luo. *SIGIR* 2023.
- [41] **HOFA: Twitter Bot Detection with Homophily-Oriented Augmentation and Frequency Adaptive Attention.**
Sen Ye, **Zhaoxuan Tan**, Zhenyu Lei, Ruijie He, Hongrui Wang, Qinghua Zheng, Minnan Luo. *arXiv preprint* 2023.
- [42] **PAR: Political Actor Representation Learning with Social Context and Expert Knowledge.**
Shangbin Feng, **Zhaoxuan Tan**, Zilong Chen, Peisheng Yu, Qinghua Zheng, Xiaojun Chang, Minnan Luo. *EMNLP* 2022.
- [43] **Heterogeneity-Aware Twitter Bot Detection with Relational Graph Transformers.**
Shangbin Feng, **Zhaoxuan Tan**, Rui Li, Minnan Luo. *AAAI* 2022.

Honors, Awards, and Research Grants

Tinker Research Grant, Thinking Machines Lab (\$5,000 in Tinker API credits)	2026
Lambda Research Grant, Lambda (\$5,000 in compute credits)	2026
OpenAI Researcher Access Program, OpenAI (\$10,000 in API credits)	2025
Top Reviewer, NeurIPS	2024
Excellent Bachelor's Thesis (ranked 1/172), Xi'an Jiaotong University	2023
AAAI Student Scholarship	2022
National Second Prize, China Undergraduate Mathematical Contest in Modeling (CUMCM)	2021
Dean's List, Xi'an Jiaotong University	2020 – 2022

Professional Service

- **Senior Program Committee Member:** AAAI 2027.
- **Area Chair:** ACL Rolling Review (2025–present).
- **Conference Reviewer:** NeurIPS (2023–; Datasets & Benchmarks 2022), ICLR (2024–), ICML (2025–), ACL Rolling Review (2023–), COLM (2024–), KDD (2024–), AISTATS (2025–), The Web Conference (2024–2025), ICWSM (2023–), Learning on Graphs (2022–2024).
- **Journal Reviewer:** TMLR, IEEE TKDE, IEEE TNNLS, IEEE/ACM Transactions on Networking.
- **Workshop Reviewer:** ACL SRW 2025; ICLR AGI 2024; EMNLP WiNLP 2024; ACL KnowledgeNLP 2024; AAAI GCLR 2024; NeurIPS Temporal Graph Learning 2023.
- **Volunteer:** EMNLP 2022, 2023.

Mentorship and Leadership

Director, LUD Lab, Xi'an Jiaotong University 2022 – 2023

Led a **17-person undergraduate research group**.

Selected Mentees (research contributions and subsequent positions, where applicable):

- **Heng Wang:** NLGraph (NeurIPS 2023 Spotlight), MVSD (EMNLP 2023) → Ph.D., UIUC
- **Yuhan Liu:** BotMoE (SIGIR 2023) → Ph.D., NYU
- **Yuyang Bai:** KGQuiz (WWW 2024 Oral) → Ph.D., Texas A&M
- **Zijian Cai:** LMBot (WSDM 2024) → M.S., Institute of Automation, CAS
- **Sen Ye:** HOFA → Ph.D., Peking University
- **Thomas Mitchell:** undergraduate researcher, University of Notre Dame

Teaching

University of Notre Dame

Teaching Assistant, CSE 40113: Design and Analysis of Algorithms	Spring 2024
Teaching Assistant, CSE 30124: Introduction to Artificial Intelligence	Fall 2023